

Canvas: Data Wrangling II on Academic Performance Dataset

I. Dataset Creation: Simulating Real-World Data Issues

Code Head: `import pandas as pd`

Code Tail: `df = pd.DataFrame(data)`

Explanation:

We start by creating a simulated dataset named **Academic performance** for 100 students. It includes these columns:

- **Student_ID**: Unique ID for each student.
- **Math, Physics, English**: Subject scores.
- **Attendance**: Value between 0.7 to 1.0 (like a percentage).
- **GPA**: Grade Point Average between 2.5 and 4.

We also introduce **intentional problems** in the dataset to mimic real-world data:

- One missing value in **Math**.
- An unrealistic **Physics** score (200).
- A letter 'A' instead of a number in the **English** score.

II. Finding Missing and Incorrect Values

Code Head: `missing_values = df.isnull().sum()`

Code Tail: `print("Inconsistencies in English_Score:",
inconsistent_data)`

Explanation:

We check each column for:

- **Missing values** using `.isnull().sum()`.
- **Invalid data types**, especially in the `English` column, by trying to convert all values to numbers using `pd.to_numeric()` and counting how many fail.

This helps us identify where data cleaning is required.

III. Visualizing the Data: Spotting Outliers

Code Head: `plt.figure(figsize=(12, 6))`

Code Tail: `plt.title('Boxplot of Numeric Variables')`

Explanation:

We use a **boxplot** to visually find outliers in numeric columns. An outlier will show as a dot or line far away from the box. Here, the Physics score of **200** will clearly appear as an outlier because valid scores should be between 0–100.

IV. Cleaning the Data: Fixing Errors and Filling Gaps

Code Head:

```
df['Math'] = df['Math'].fillna(df['Math'].mean()).astype(float)
```

```
df['Physics'] = df['Physics'].apply(lambda x: np.nan if x > 100 else x)
```

```
df['English'] = pd.to_numeric(df['English'], errors='coerce')
```

Code Tail: `df.dropna(inplace=True)`

Explanation:

We fix the issues found earlier:

- **Math:** Fill missing value using the column's average score.

- **Physics:** Change any score > 100 (like 200) to NaN because it's invalid.
- **English:** Convert all values to numbers; 'A' becomes NaN automatically.

Finally, we **drop all rows** that still have any missing values after cleaning.

V. Second Pass: Detecting More Subtle Outliers

Code Head: `z_scores = np.abs(zscore(df[numeric_cols]))`

Code Tail: `outlier_indices = np.where(z_scores > threshold)[0]`

Explanation:

We now apply a **statistical method** using the **Z-score**:

- Z-score tells how far a value is from the mean.
- We define an outlier as any value with a Z-score above 3 (i.e., it's very far from normal).
- We find all row indices where this happens.

This helps find **hidden outliers** that weren't obvious in the boxplot.

VI. Final Cleanup: Removing Outliers and Transforming Data

Code Head:

```
df_cleaned = df.drop(outlier_indices)
```

```
df_cleaned['GPA'] = np.log1p(df_cleaned['GPA'])
```

Code Tail:

```
sns.histplot(df_cleaned['GPA'], kde=True)
```

```
plt.title('Transformed GPA Distribution')
```

```
plt.show()
```

Explanation:

- We remove all rows that had outliers based on Z-scores.
- We apply **log1p transformation** ($\log(x + 1)$) to the **GPA** column.
 - This makes the GPA values less skewed (i.e., more normally distributed).
 - It helps when you're using machine learning or statistics later.

Finally, we show a **histogram** of the transformed GPA to see its new shape.

VII. Summary and Justification

Code Head: *(Conceptual section)*

Code Tail: *(Conceptual section)*

Explanation:

Why we did what we did:

- **Missing values** were filled (mean for Math) or removed (for other columns).
- **Incorrect values** (like 'A', or 200) were fixed or removed.
- **Outliers** were handled using a statistical method (Z-score).
- **Data transformation** helped in normalizing GPA, useful for further modeling or analysis.

This makes the dataset **clean, consistent, and ready for analysis or machine learning**.