

1. What is Logistic Regression?

Logistic Regression is a statistical method for binary classification. It predicts the probability that an input belongs to a certain class using the logistic (sigmoid) function.

2. What kind of problem does this code solve?

This code solves a binary classification problem: predicting whether a user purchases a product or not based on age and estimated salary.

3. What is the sigmoid function?

The sigmoid function maps input values to a range between 0 and 1. It helps interpret model outputs as probabilities in logistic regression.

4. Why did we use StandardScaler?

We use **StandardScaler** to normalize the feature values so they have a mean of 0 and a standard deviation of 1. This improves model convergence and accuracy.

5. What is a confusion matrix?

A confusion matrix is a 2x2 table for binary classification problems showing true positives (TP), false positives (FP), true negatives (TN), and false negatives (FN).

6. What do TP, FP, TN, FN mean?

- TP: Correctly predicted positives
- FP: Incorrectly predicted positives

- **TN: Correctly predicted negatives**
 - **FN: Incorrectly predicted negatives**
-

7. What is accuracy?

Accuracy is the ratio of correctly predicted samples (TP + TN) to total predictions. It shows how often the classifier is correct.

8. How do you calculate precision?

Precision = $TP / (TP + FP)$

It indicates how many of the predicted positives are actually positive.

9. How do you calculate recall (sensitivity)?

Recall = $TP / (TP + FN)$

It measures how many actual positives were correctly identified.

10. What is the error rate?

Error Rate = $(FP + FN) / \text{Total predictions}$

It represents the proportion of incorrect predictions.

11. What is specificity?

Specificity = $TN / (TN + FP)$

It measures how well the model identifies negatives.

12. Why is logistic regression suitable here?

It is suitable because the target variable is binary — purchase (1) or not (0) — which is ideal for logistic regression.

13. What are the independent and dependent variables?

Independent: Age and Estimated Salary

Dependent: Purchased (0 or 1)

14. What does the `.fit()` function do?

`.fit()` trains the logistic regression model using the training data.

15. Why did we use `train_test_split()`?

To divide the dataset into training and test sets, ensuring the model is trained and evaluated on different data for unbiased performance.

16. What is the role of `random_state=42`?

It ensures reproducibility of the train-test split by setting a seed for the random number generator.

17. What evaluation metrics did we use?

Accuracy, Precision, Recall, Error Rate, Specificity, F1-Score, and Confusion Matrix were used for evaluation.

18. What does `.predict()` do?

`.predict()` uses the trained model to predict the class labels for the test data.

19. What is F1-score?

The F1-score is the harmonic mean of precision and recall. It balances both metrics in one score.

20. What does `.transform()` do in StandardScaler?

It applies the learned scaling (mean and std) to the dataset, transforming it to a standardized format.

21. Why is scaling important in logistic regression?

Because logistic regression is sensitive to the scale of features. Unscaled features can dominate others and affect predictions.

22. What solver did we use in Logistic Regression?

We used `lbfgs`, a robust solver suitable for small to medium datasets.

23. How does Logistic Regression handle non-linearity?

Logistic regression assumes a linear relationship between features and the log-odds. It doesn't handle non-linearity well without feature engineering.

24. What is the output of a logistic regression model?

The model outputs a probability score between 0 and 1, which is then thresholded (usually at 0.5) to classify into 0 or 1.

25. Why do we visualize the confusion matrix?

To quickly see how many predictions were correct vs. incorrect in each category, helping us assess model performance clearly.